Using AI at Mind

AI at Mind powers more than our robots. It powers how our team works and delivers on our mission.

Everyone at Mind should use AI, in nearly every aspect of our work. But as importantly, Mind should be the place where, decades from now, you can say you learned how to use AI *skillfully*.

Because like any technology, AI can be used poorly or brilliantly.

This doc exists to guide how we use AI at Mind. It's a living doc and we'll update it regularly. You can comment on it directly with suggestions and questions.

What this is not: Tips for how to use AI for specific use cases, e.g. evals, financial analysis, recruiting, software development. Expect those golden paths to be published over time by the teams closest to each type of work.

These are Mind's first principles for skillful AI use:

You are responsible for anything you put your name on. How you made it is your business. We're not asking anyone to disclose that they used AI, or to label documents, or to caveat their work, because the tool involved isn't the interesting thing about your work. Whether the work is right and whether it's clear...that's our accountability bar. What we do expect is that when someone asks you about something you delivered, you have an answer that reflects that you truly understand the content. This applies to slides, memos, analyses, Slack messages, customer emails, and code.

Read your own work before you ask anyone else to. A reasonable AI use test is whether you could talk someone through the document without looking at it, and answer questions they might have about any part of it. If you can't, what you're really doing is handing your teammate the review that you skipped. There's an easy out here that we'd like people to use more often: Just say what you don't know. Sending something with "I generated most of this and I'm shaky on the middle section, can you look hard at it?" is completely fine, and usually faster for everyone than pretending otherwise.

It should take you longer to produce something than it takes someone else to consume it. Most of what we produce gets built once and read by a lot of people, so any slop you leave in your work gets

Use the time you save with AI to think more, not less. The gain with AI isn't really that you can produce the same work faster; it's that you can afford to consider four approaches where you used to have time for one. Ask the model to argue against your conclusion. Ask what a skeptical reader would push on. Ask it to find the weakest part of your draft, then decide for yourself whether it's right. If you get all the way through a task without changing your mind about anything, that's usually a sign you were transcribing a preconceived conclusion rather than deeply thinking.

Check things in proportion to what happens if they're wrong. Brainstorming needs little to no verification. On the other hand, confirm data that's going to influence someone else's decision. AI is often most convincing in the places it's least reliable. There's no general rule for how much checking is enough, which is why this is a principle rather than a process. Make the call based on the stakes.

Work with AI in public where you can. Working in public is rocket fuel for Mind's rate of learning. So if you're about to build something with AI, say so in #ai-help before you start rather than after. Share the prompts and workflows that worked, including the unglamorous ones. [AI can help you with this.] Most of the leverage available to us right now sits in other people's best practices rather than in the tools themselves. Good habits spread when people share them. And with help from AI, we can convert individuals' best ways of working into team-level standards.